

C++

February 5, 2015

Mike Spertus

mike_spertus@symantec.com





Matrix example

- For the rest of the quarter, we are going to create a linear algebra library to see how to combine the C++ features we have been learning to make an awesome library
- We are (very) loosely inspired by the matrix classes from the Origin C++ libraries



Matrices

- A matrix is just a two dimensional array of numbers (picture from Wikipedia)

$$\begin{bmatrix} 1 & 9 & -13 \\ 20 & 5 & -6 \end{bmatrix}.$$

- Matrices are used in all branches of science, statistics, economics, math, etc. and can be added, multiplied, or have their determinants taken
- I will give you all necessary formulas



Initializing matrices

- We would like our Matrix class to have a natural initializer like the following 2x3 matrix
- `Matrix<2,3> m = { {2, 4, 6},
 {1, 3, 5}};`

Another kind of constructor

Initializer lists



- In C++98, one of the problems with using vectors was that it was difficult to initialize them

- `// Yuck!`
`int init[] = { 0, 1, 1, 2, 3, 5};`
`vector<int> v(init,`
`init+sizeof(init)/sizeof(init[0]));`

Can't user-defined types be as easy to init as built-in ones?



- In C++11, you can initialize vectors as easily as C arrays
 - `vector<int> v = {0, 1, 1, 2, 3, 5};`
- How does `vector<T>` do this?
- It has a constructor that takes a `std::initializer_list<T>`, which represents a “braced initializer of Ts” expression
- Initializer lists have `begin()`, `end()`, and `size()` methods so your constructor can iterate through their value.

Initializer list constructor for matrix



- The code to initialize as mentioned previously is on chalk
- Note from the slide above that we initialize matrices with “initializer lists of initializer lists”
 - The initializer list contains an initializer list for each row
- ```
Matrix(initializer_list<initializer_list<double>> init) {
 auto dp = data.begin();
 for (auto row : init) {
 std::copy(row.begin(), row.end(), dp->begin());
 dp++;
 }
}
```



# Matrix arithmetic

- In order to easily add and multiply matrices, we would like to be able to tell “+” and “\*” about matrices
- This is called operator overloading





# Operator overloading

- You can overload operators just like functions
- The following operators can be overloaded:
- Unary operators:

**+ - \* & ~ ! ++ -- -> ->\***

- Binary operators:

**+ - \* / % ^ & | << >>  
+= -= \*= /= %= ^= &= |= <<= >>=  
< <= > >= == != && ||  
, [] ()  
new new[] delete delete[]**

# Which operators can't be overloaded?



- `., .* , ? : , ::`
- Fame and fortune await for the one who figures out how to overload “operator.()”

# Operator overloading examples



```
class myString {
 myString(const char *cp);
 char operator[](size_t idx) const;
 myString operator+(myString &addend) const;
 myString operator+=(myString &addend);
 inline friend myString
 operator+(const myString &s1, const myString &s2) {
 }
};

// Alternatively
myString
operator+(const myString &s1, const myString &s2);
```

# Which way of overloading addition is better?



- Consider `"Hello " + myString("World")`
- Doesn't work for the member function
  - The first argument isn't even a class, so the compiler wouldn't know where to look for a member function.
- What about `myString("World") + "Hello "`
  - Works for both
- Using a global function makes sure both arguments are treated the same way, which fits the intuition that addition operators, which are generally commutative, should apply the same rules to each arguments.



# Do the same way for printing

- ostream &  
operator<<(ostream &os,  
                    myString const &ms) {...}
- If you want to be fancy  
template<typename charT, typename traits>  
basic\_ostream<charT, traits>  
&operator<<(basic\_ostream<charT, traits>  
&os, myString const &ms) {...}

# How does an I/O manipulator get invoked



- Recall that `endl` is defined (as modulo some template complication that is irrelevant here) follows
- ```
ostream &
endl(ostream &os)
{
    os << '\n';
    os.flush();
    return os;
}
```
- How come “`cout << endl;`” actually behaves as “`endl(cout)`”?

Another overload!



```
ostream &
operator<<
(ostream&os,
 ostream&(*manip)(ostream &))
{
    return manip(os);
}
```

How does a smart pointer work?



- Overloading operator->() of course
- operator->() overloads with a unique rule
 - Keep doing -> until it is illegal



shared_ptr

- shared_ptr is a reference counted pointer.
- Inside shared_ptr, we have something like:

```
template<class T>
class shared_ptr {
    // Returns the wrapped pointer
    T *operator->();
};
```
- Because the -> is applied again, it acts just like the wrapped pointer except that it maintains a reference count.

Flow of control is hard to follow but memory is easy to manage



```
class A;  
int f()  
{  
    shared_ptr<A> ap1 (new A());  
    shared_ptr<A> ap2 (new A());  
    return ap1->i + ap2->i;  
}
```

- Deletes automatically no matter what



Example

```
#include <boost/shared_ptr.hpp>

int f()
{
    auto ap = make_shared<A>();
    ap->m(1);
    g(ap);
} // the A object is automatically deleted
// when we leave scope unless someone
// else is using it
```



Best practice

- All dynamic duration objects should be owned by a smart pointer
- Not all uses need to be through a smart pointer, but the owner needs to be one



Overloading operator++()

- To overload ++x, write
`X &X::operator++() { ... }`
- To overload x++, write
`X &X::operator++(int) {...}`
- The int argument isn't really there. Don't use it! The signature just gives a way to distinguish preincrement and postincrement



Overload && and ||

- The built-in operator && (logical and) has “short circuit evaluation”
 - If the left argument is false, the right one isn’t evaluated because it can’t make the && true.
 - `i != 0 && 5/i > 0` // Will never divide by 0
- User-defined overloads of && and || never short circuit. Both arguments are evaluated no matter what.

Overloading operators for matrices



- The sample Matrix.h on chalk overloads matrix multiplication with the (complicated) rule for multiplying matrices



Specializing templates

- The “secret sauce” for C++ templates is that if the general “generic” definition of the template isn’t really what you want for a particular set of template parameters, you can override it for that particular case with a specialization
- Think of this as the compile-time analog to object orientation where you also override a more general method in a more specialized derived class



Matrix determinants

- The determinant is a number that represents “how much a matrix transformation expands its input”
 - Don't worry if you don't understand this
- We will just use the formula to calculate them
- The general formula is here
 - http://en.wikipedia.org/wiki/Laplace_expansion
- But I will give you the special case you need for the homework



Full specialization

- A function, class, or member can be fully specialized
- See the definition of `Matrix<1,1>::determinant()` in `Matrix.h`



Partial specialization

- Only classes may be partially specialized
- Template class:

```
template<class T, class U>  
class Foo { ... };
```
- Partial specialization:

```
template<class T>  
class Foo<T, int> { ... };
```
- You can tell the second is a specialization because of the `<>` after the class name



Partial specialization

- The partially specialized class has no particular relation to the general template class
 - In particular, you need to either redefine (bad) or inherit (good) common functionality
 - For example, see PSMatrix.h



Overloading

- functions cannot be partially specialized, so overloading is used instead
- For example, see `OverloadMatrix.h`

Compilation of template methods



- A method of a template class is only compiled if it is used
 - Indeed this is true for any kind of template
- That's why in the Matrix example `Matrix<int, 1, 1>` objects can be instantiated even though `Matrix<int, 1, 1>::minor(int, int)` doesn't compile
- This has interesting implications for static members



Exercise 5-1

- Modify Matrix.h to let you add matrices
- To add two matrices, they both have to have the same number of rows and columns
- Just add the corresponding elements to get a new matrix with the same number of rows and columns
- See <http://www.purplemath.com/modules/mtrxadd.htm> for an example



Exercise 5-2

- For each of the following programs, modify them to have a direct (i.e., specialized or overloaded implementation) of determinants for 2x2 matrices.
 - Matrix
 - PSMatrix
 - OverloadMatrix
- The formula for the determinant of the 2x2 matrix m is
$$m(0, 0) * m(1, 1) - m(1, 0) * m(0, 1)$$
- Test how much your code changed the execution time for the programs. What do you conclude?



Exercise 5-3: Extra credit

- Some of the calculations in determinant seem to copy matrices a lot
- Can you modify the matrix class to be (efficiently) movable
- Does that improve the benchmark?



Exercise 5-4

- Create a `ComplexInt` class that acts like a complex integer (`c.r` is the real part. `c.i` is the imaginary part)
 - Define multiplication and addition for complex integers
 - Ensure that `cout << c;` prints something like `5+3i`.



Exercise 5-5: Extra credit

- Look up how to create user-defined literals
- Create a user-defined literal to make it easier to enter complex numbers

```
ComplexInt ci = 6 + 3_i;
```



Exercise 5-6: Extra Credit

- What happens when you do

```
ComplexInt ci(2, 7); // 2 + 7i  
cout << setw(10) << ci << endl;
```
- Is this the desired behavior?
- If not, how would you fix it?