

C++

February 19, 2015

Mike Spertus

mike_spertus@symantec.com





Lambdas

- We have used C++11 lambdas in class
- `[] (int x, int y) { return x < y; }` is an anonymous boolean valued “function” returning true if $x < y$
- “Function” is in quotes because it is actually an unspecified type. You can only assign to a function pointer if there is no capture list (see following slides), but you can always store as follows
 - `auto f = [] (int x, int y) { return x < y; }`



Lambda return values

- If your lambda just consists of a return statement, the compiler infers the type.
- If not, you give it using the “unified function syntax”

```
[ ] (int x, int y) -> double {  
    int z = x + y; return z + x;  
}
```

- This syntax avoids parsing problems



Capture lists

- To capture local variables by reference, use [&]
- To capture local variables by value, use [=]
- For finer-grained results, specify exactly what you want to capture

```
• int i=0;  
  vector<int> v;  
  /* ... */  
  for_each(v.begin(), v.end(),  
           [&i](int x) { i += x; });
```



Lambdas and algorithms

- Now all of the standard library algorithms can be used as easy as for loops
- ```
std::vector<int> someList; // Wikipedia
int total = 0;
std::for_each
 (someList.begin(), someList.end(),
 [&](int x) { total += x; });
```

# C++14: Polymorphic lambdas



- In C++14, you can give a lambda an `auto` argument
- In the following code, the compiler figures out that `x` is an `int`.
- ```
std::vector<int> someList; // Wikipedia  
int total = 0;  
std::for_each  
    (someList.begin(), someList.end(),  
     [&](auto x) { total += x; });
```

Behind the scenes: How polymorphic lambdas work



- Advanced
- When the compiler sees
 - `[] (auto x) { cout << x; }`
- It generates a functor
 - ```
struct lambdaFunctor_42393626243 {
 template<typename T>
 void operator() (T t) {
 cout << t;
 }
};
```



# C++11 THREADS





# Overview

- Perhaps the biggest addition to C++11 is support for standardized concurrency
  - Multithreading to run tasks in a process in parallel with each other
  - Synchronization primitives and memory model to allow different threads to safely work with the same data



# Why is this a big deal?

- Perhaps the biggest secret in computer progress is that computer cores have not gotten any faster in 10 years
  - 2005's Pentium 4 HT 571 ran at 3.8GHz, which is better than many high-end CPUs today
  - The problem with increasing clock speeds is heat
    - A high end CPU dissipates over 100 watts in about 1 cubic centimeter
    - A light bulb dissipates 100 watts in about 75 cubic centimeters

# Why doesn't anyone know about this?



- Even though cores have not gotten faster, the continued progression of Moore's law means that computers today have many cores to run computations in parallel
  - Even cell phones can have 4 cores
  - 12 to 24 cores are not unusual on high-end workstations and servers
    - 24 to 48 if you count hyperthreading



# Back to C++

- Unfortunately, C++ did not have any notion of multithreading until C++11 came out
- C++ programmers used os-provided multithreading libraries like pthreads and win32 threads
- But this is not acceptable
  - Using these libraries are clunky, not well integrated with other language constructs, and not C++ like
  - Even worse, Threads Cannot be Implemented as a Library (Hans Boehm, PLDI 2005)
    - <http://www.hpl.hp.com/techreports/2004/HPL-2004-209.pdf>



# References

- C++ Concurrency in Action Book
  - <http://www.manning.com/williams/>
    - If you buy from Manning rather than Amazon, you can download a preprint right now without waiting for the official publication
  - The author Anthony Williams is one of the lead architects of C++11 threads, the maintainer of Boost::Thread, and the author of just::thread
- Anthony's Multithreading in C++0x blog
  - <http://www.justsoftwaresolutions.co.uk/threading/multithreading-in-c++0x-part-1-starting-threads.html>
  - Free with concise coverage of all the main constructs
- The standard, of course
  - Also look at the papers on the WG21 [site](#)



# WARNING!

- The next several slides are very confusing
- You do not need to learn them in detail (or at all) as C++11 resolves these problems
- However, we give these slides for two reasons
  - Without seeing such bizarre unexpected behavior, one would be tempted to continue using thread libraries
  - They are very interesting

# What can r1 and r2 end up as? (Boehm)



Initially  $x = y = 0$ ;

## Thread 1

```
x = 1;
r1 = y;
```

## Thread 2

```
y = 1;
r2 = x;
```

# Answer: Any combinations of 0 and 1!



- Intuitively  $r1 == r2 == 0$  impossible
- Practically, the compiler (or the hardware) may reorder the statements because it doesn't matter within a given thread which order the assignments take place
- However, it does matter if the variables are used by another thread at the same time and we could end up with both  $r1$  and  $r2$  being 0
  - Note: Under pthreads rules this is simply illegal





# If $q = 0$ , what can another thread see count as? (Boehm)

```
[count is global]
```

```
for (p = q; p != 0; p = p->next) {
 count++;
}
```

- Other threads may see `count == 1`!
- Compiler may rewrite code by speculatively incrementing `count` before the loop, and decrementing if necessary at the end!
- Even `gcc -O2` does this.

# Is this code correct?



```
class A {
public:
 virtual void f();
};
A *a; // Global variable
```

## Thread 1

```
a = new A;
```

## Thread 2

```
if (a) a->f();
```

# Not on modern multicore computers!



- Writes made on one processor may not be seen in the same order on another processor!
  - Allows microprocessor designers to use write buffers, instruction execution overlap, out-of-order memory accesses, lockup-free caches, etc.
- Thread 2 may see the assignment to `a` before it sees the vtable of the new `A` object!
- If that happens, the `a->f()` call will crash!
- Modern processors use *Weak Consistency*

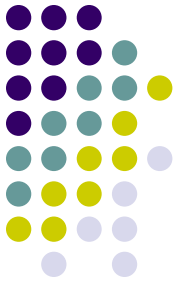
# Weak memory consistency



In a multiprocessor system, storage accesses are weakly ordered if (1) accesses to global synchronizing variables are strongly ordered, (2) no access to a synchronizing variable is issued by a processor before all previous global data accesses have been globally performed, and if (3) no access to global data is issued by a processor before a previous access to a synchronizing variable has been performed.

—Dubois, Scheurich, Briggs (1986)

If the compiler does not have a notion of synchronizing variables, the above says nothing! Prior to C++11, this is addressed non-portably by vendor-specific synchronization extensions to C++.



# THE BASICS



# Hello, threads

```
#include <iostream>
#include <thread>

void hello_threads() {
 std::cout<<"Hello Concurrent World\n";
}

int
main() {
 // Print in a different thread
 std::thread t(hello_threads);
 t.join(); // Wait for that thread to complete
}
```



# What happened?

- Constructing an object of type `std::thread` immediately launches a new thread, running the function given as a constructor argument (in this case, `hello_threads`).
  - We'll talk about passing arguments to the thread function in a bit.
- Joining on the thread, waits until the thread completes
  - Be sure to join all of your threads before ending the program
  - Exception: Later we will discuss detached threads, which don't need to be joined



# Locks

- The simplest way to protect shared data is with a `std::mutex`.
- How can we make sure we release the mutex when we are done no matter what?
- RAI!
- C++11 includes a handy RAI class `std::lock_guard` for just this purpose.



# Locks



```
#include <list>
#include <mutex>
#include <algorithm>
```

```
std::list<int> some_list; // A data structure accessed by multiple threads
std::mutex some_mutex; // This lock will prevent concurrent access to the shared data structure
```

```
void
add_to_list(int new_value)
{
 std::lock_guard<std::mutex> guard(some_mutex); // Since I am going to access the shared data struct, acquire the lock
 some_list.push_back(new_value); // Now it is safe to use some_list. RAll automatically releases lock at end of function
}
```

```
bool
list_contains(int value_to_find)
{
 std::lock_guard<std::mutex> guard(some_mutex); // Must get lock every time I access some_list
 return
 std::find
 (some_list.begin(),some_list.end(),value_to_find)
 != some_list.end();
}
```

# Not so basic: Thread arguments



- You can add arguments to be passed to the new thread when you construct the `std::thread` object as in the next slide
- But there are some surprising and important gotchas that make passing arguments to thread function different from passing arguments to ordinary functions, so read on

# Passing arguments to a thread



```
#include <iostream>
#include <thread>
#include <string>
#include <vector>
#include <mutex>
using namespace std;
mutex io_mutex;

void hello(string name) {
 lock_guard<mutex> guard(io_mutex);
 cout <<"Hello, " << name << endl;
}

int
main(){ // No parens after thread function name:
 vector<string> names = { "John", "Paul"};
 vector<thread> threads;
 for(auto it = names.begin(), it != names.end(); it++) {
 threads.push_back(thread(hello, *it));
 }
 for(auto it = threads.begin(), it != threads.end(); it++) {
 it->join();
 }
}
```



# Deceptively simple

- A different notation is used from arbitrary function calls, but otherwise fairly straightforward looking:

- ```
void f(int i);  
f(7); // Ordinary call  
thread(f, 7); // f used as a thread function
```

Gotcha: Passing pointers and references



- Be very careful about passing pointers or references to local variables into thread functions unless you are sure the local variables won't go away during thread execution
- Example (based on Boehm)

```
void f() {  
    int i;  
    thread t(h, &i);  
    bar(); // What if bar throws an exception?  
    t.join(); // This join is skipped  
} // h keeps running with a pointer  
    // to a variable that no longer exists  
    // Undefined (but certainly bad) behavior
```
- Use try/catch or better yet, a RAII class that joins like the `thread_guard` class in *Concurrency In Action* book

Gotcha: Signatures of thread functions silently “change”



- What does the following print?

```
void f(int &i) { i = 5; }  
int main() {  
    int i = 2;  
    std::thread t(f, i);  
    t.join();  
    cout << i << endl;  
    return 0;  
}
```

A compile error (if you're lucky), 2 if your not!



- Of course, 5 was intended
- Unfortunately, thread arguments are not interpreted exactly the same way as just calling the thread function with the same arguments
- This means that even an application programmer using threads needs to understand something subtle about templates



What went wrong, continued

- Imagine `std::thread`'s constructor looks like the following

```
struct thread { ...  
    // 0 arg thrfunc constructor  
    template<typename func>  
    thread(func f);  
    // 1 arg thrfunc constructor  
    template<typename func, typename arg>  
    thread(func f, arg a);  
    ...  
};  
...  
    // Deduces thread::thread<void(*) (int), int>  
    std::thread t(f, i);  
...  
};
```

- In fact, thread constructors use “variadic argument lists,” which we haven’t (yet) covered

IOW, Templates don't know f takes its argument by reference



- To do this, we will use the “ref” wrapper in `<functional>`
- ```
void f(int &i) { i = 5; }
int main() {
 int i = 2;
 std::thread t(f, std::ref(i));
 t.join();
 cout << i << endl;
 return 0;
}
```

# Does thread's constructor really look like that?



- No, C++11 has “variadic templates” that can take any number of arguments, so we don’t need to separate 0-arg, 1-arg, etc. constructors:

```
struct thread {
 template
 <typename F, typename... argtypes>
 thread(F f, argtypes... a);
 ...};
```
- We’ll learn about these later



# Lambdas can help

- As a result of the above complexity, some people recommend using lambdas instead of functions or functors with threads
- ```
int main() {  
    int i = 2;  
    std::thread t([&i]() {i=5;});  
    t.join();  
    cout << i << endl;  
    return 0;  
}
```
- Correctly prints 5



Thread local storage

- A new storage duration.
- Each thread gets its own copy
- `thread_local int i;`

Async functions: Running functions in another thread



- It's nice that we can pass arguments to a thread (like we do to functions), but how can we get the thread to return a value back?
- Basically, we want to be able to use threads as “asynchronous functions”
- C++11 defines a `std::future` class that lets a thread return a value when it's done
- Create a future with `std::async`
 - As soon as you create it, it starts running the function you passed it in a new thread
 - Call `get()` when you want to get the value produced by the function
 - `get()` will wait for the thread function to finish, then return the value
 - See example below



std::future example

- From [Multithreading in C++0x Part 8](#)

```
#include <future>
#include <iostream>

int calculate_the_answer_to_LtUaE();
void do_stuff();

int main()
{
    std::future<int> the_answer
        = std::async(calculate_the_answer_to_LtUaE);
    do_stuff();
    std::cout <<"The answer to life, the universe and everything is "
              << the_answer.get()
              << std::endl;
}
```

Can I check if the future has a value yet?



- Yep, `std::future` has an `is_ready()` method that tells you if the thread function has completed.

What if the asynchronous function throws an exception?




- If the thread function in a future throws an exception instead of returning a value, then calling `get()` will throw the exception, just like the asynchronous function was a real function



Parallel accumulate

- It would be really nice to have an implementation of `std::accumulate` that breaks up its input into pieces, adds up each piece in parallel and then adds up the results from each of the pieces
- Let's do this with futures
- `async_accumulate_function.cpp`

Homework





HW 7-1

- The purpose of this problem is to ensure that you can write basic multithreaded code on your system. Since threading is not portable, please send a transcript. Use the C++11 threads
- Write a program that creates 3 threads that each count up to 100 and output lines like:
Thread 3 has been called 4 times
- To get a thread number, use
`std::this_thread.get_id()`
- Make sure you use synchronization to keep different threads from garbling lines like the above.
- Submit the output from your program. What does it tell you about how threads are actually scheduled on your system?



HW 7-2

- Write a thread-safe stack using locks
 - The only required operations are push and pop
- E.g., multiple threads can concurrently do things like the following without corrupting the stack
 - ```
mpcs50144::stack<int> s;
s.push(7);
s.push(5);
cout << s.pop();
```
- For extra credit, add additional useful functionality (e.g., initializer list constructor, etc.)



## HW 7-3

- Modify the matrix program to compute the determinant in parallel
- How does the performance depend on the size of the matrix?
- For extra credit, come up with other improvements to the performance of determinant



## HW 7-4

- Use a lambda with a capture list to give another solution to the `for_each` homework from lecture 4 that is different from all of the solutions I gave in class during the week 5 homework review
  - I have uploaded them to this week's chalk
- Compare and contrast it to the other solutions