C++
March 5, 2015

Mike Spertus

mike_spertus@symantec.com

Pointers

- Pointers to a type contain the address of an object of the given type or inheriting from the given type.
  ```
  A *ap = new A;
  ```
- Arithmetic on pointers
  ```
  A *aap = new A[10];
  ```
  *(aap + 5) is the 5th element of the array
  It doesn't add 5 to the address, but adds enough to get to the fifth element (starting from 0)
- Dereference with `*`
  ```
  A a = *ap;
  ```
- `->` is an abbreviation for `(*_)`.
  ```
  ap->foo(); // Same as (*ap).foo()
  ```
- If a pointer is not pointing to any object, you should make sure it is nullptr
  ```
  ap = nullptr; // don't point at anything
  if(ap) { ap->foo(); }
  ```

Smart pointers that own their data

- unique_ptr<T> points to a T object. Its destructor deletes the object

- Is the following correct?

  - unique_ptr<char> = new char[10];

Unique pointers to arrays

- No!
- We need to delete the array with delete[], but unique_ptr<char> uses delete
- With ordinary pointers, we manually remember it's pointing to an array and then cause the correct deletion function
- How can we tell the unique_ptr to automatically use the correct operator delete?
- Answer: Use a unique_ptr to an array
  - `unique_ptr<char[]> = new char[10];`

How to transfer ownership into an out of a unique_ptr

- ```
  unique_ptr<A> ap(new A); // ap owns the new A
  // delete old A and own new A
  ap.reset(new A);
  ```
- **Suppose we have the following functions**
- ```
  void f1(A *); // f1 will not del
  void f2(A *); // f2 is responsible for deleting
  // Funcs returning new obj should
  // return a unique_ptr
  unique_ptr<A> Afactory();
  f1(ap.get()); // Pass the pointer
  // Pass ptr and release ownership
  f2(ap.release());
  ap = Afactory();
  // You can't copy unique pointers because that
  // would create ambiguous ownership. You
  // move them instead
  unique_ptr<A> ap2 = move(ap);
  ```

make_unique

- Since using owning raw pointers is error-prone, we prefer not to ever manually call new and delete
- Instead, we want creation functions that return owning smart pointers
- make_unique creates a new object and returns a rvalue unique_ptr
- ```
  auto ap = make_unique<A>(1, 2, 3);
  ```
  is the same as
  ```
  unique_ptr<A> ap(new A(1, 2, 3));
  ```
  but you are never exposed to new, delete or raw pointers
- Best practice: Always use an owning pointer

shared_ptr

- Just like unique_ptr only can have multiple pointers to same object
- Once all of the pointers go away, the object is deleted
- Can create with make_shared
- A difference from unique_ptr
  - Can copy (That's how you get multiple owners)
- Don't create two shared_ptr's from the same object
  - You'll get two reference counts!
  - Advanced: Read about enable_shared_from_this if you run into a situation where you need to create a shared_ptr from the this pointer

Pointers to functions

- The basic idea is usually that you describe a type by how it is used
  - `int *ip; // Means *ip is an int`
  - `int (*fp)(int, int); // *fp can be called with 2 ints`
- Let's show fp in action
  - ```
    int f(int i,int j) { … }
    fp = &f;
    fp(2, 3);
    // The following line only works without captures
    fp = [](int i, int j) { return i + j; }
    ```

Function pointers: motivation

- Sometimes we don't know what function we want to call until runtime

- ```
  double mean(vector<double> const &) {...}
  double median(vector<double> const &) { ... }
  cout << "Should I use means or medians ";
  string answer;
  cin >> answer;
  double (*averager)(vector<double> const &)
    = (answer == "mean" ? mean : median);
  cout << "The average home price is ";
  cout << averager(getHomePrices()) << endl;
  ```

Pointers to members

- ```
struct A {
  int i;
  int j;
 void foo(double);
 void bar(double);
};
```
- We would like to be able to point to a particular member of A
  - Not an address because we haven't specified an A object
  - More like an offset into A objects
- ```
int A::*aip = &A::i;
void (A::*afp)(double) = &A::foo;
A *ap = new A;
A a;
ap->*aip = 3; // Set ap->i to 3
(a.*afp)(3.141592); // Calls a.foo(3.141592)
```

Pointer to member functions

- Consider
```
vector<Animal *> zoo;

zoo.push_back(new Elephant);
zoo.push_back(new Zebra);
zoo.push_back(new Bear);
cout << "Feeding time (f) or Bedtime (b)?"
char c;
cin >> c;
auto ap
  = c == 'f' ? &Animal::eat : &Animal::sleep;

for(auto animal : zoo) {
  animal->*ap();
}
```

Using with standard smart pointers

- Unfortunately, `unique_ptr` and `shared_ptr` don't overload `operator->*()`, so if we want to make the previous example delete objects when the zoo closes (or there is an exception when constructing an animal), we should modify it as shown below

- ```cpp
  vector<unique_ptr<Animal>> zoo;
  zoo.emplace_back(new Elephant);
  zoo.emplace_back(new Zebra);
  zoo.emplace_back(new Bear);
  cout << "Feeding time (f) or Bedtime (b)?"
  char c;
  cin >> c;
  void (Animal::*ap)()
    = c == 'f' ? &Animal::eat : &Animal::sleep;

  for (auto it = zoo.begin(); it != zoo.end(); it++) {
    ((**it).*ap)();
  }
  ```

References

- Like pointers but different
  - Allow one object to be shared among different variables
  - Can only be set on creation and never changed
    - Reference members must be initialized in initializer lists
      ```
      struct A {
        A(int &i) : j(i) {}
        int &j;
      };
      ```
  - Cannot be null

Not all callables can be assigned to a function pointer

- Can only assign a lambda to a function pointer if it does not have a capture list
  - See homework
- Can't assign a functor to a function pointer
-
```cpp
struct WeightedMean {
  WeightedMean(vector<double> const &weights)
    : weights(weights) {}
  double operator()(vector<double> const &data) {
    return
      inner_product(data.begin(), data.end(),
                    weights.begin(), 0.0)
      / accumulate(weights.begin(), weights.end(), 0.0);
  }
  vector<double> weights;
};
double (*averager)(vector<double> const &)
  = WeightedMean({1.5, 3.6, 4.2}); // Error!
cout << "The average home price is ";
cout << averager(getHomePrices()) << endl;
```

std::function

- We have just discussed function pointers, but in C++, functions aren't the only thing that can be called
  - Call a function
  - Call a lambda
  - Call a functor
  - Call a member function

std::function can hold anything callable

- ```cpp
  struct WeightedMean {
    WeightedMean(vector<double> const &weights)
      : weights(weights) {}
    double operator()(vector<double> const &data) {
      return
        inner_product(data.begin(), data.end(),
                      weights.begin(), 0.0)
        / accumulate(weights.begin(), weights.end(), 0.0);
    }
    vector<double> weights;
  };
  function<double(vector<double> const &)> averager
    = WeightedMean({1.5, 3.6, 4.2}); // OK
  cout << "The average home price is ";
  cout << averager(getHomePrices()) << endl;
  ```

You can even put a member pointer in a std::function

- It acts like a function whose first argument is the "this" pointer (or even a reference).

Often you can choose between a std::function and a template

- In the below code, `tmpl_apply` and `fn_apply` can be used similarly
- ```
  template<typename Callable>
  double
  tmpl_apply(Callable c, vector<double> const &data)
  {
    return c(data);
  }

  double
  fn_apply(function<double(vector<double> const &)> c, vector<double> const
  &data)
  {
    return c(data);
  }
  void f()
  {
    tmpl_apply(mean, {1.7, 2.3}); // OK
    fn_apply(mean, {1.7, 2.3}); // OK
    tmpl_apply(WeightedMean({1.2, 3.4}), {1.7, 2.3}) // OK
    fn_apply(WeightedMean({1.2, 3.4}), {1.7, 2.3}) // OK
  }
  ```
- See HW

Standard exception types

- Even though technically, you can throw exceptions of any type, you should always have your exceptions inherit from std::exception, std::runtime_error, or std::logic_error

- Another good best practice: Throw by value but catch by reference

- Remember, don't use exception specifications
  - But note that C++11 introduces noexcept keyword (beyond the scope of this quarter, but will come back to it next quarter)

Tuples

- Tuples are a generalization of std::pair to any number of fields

```
pair<string, int> si = make_pair("str", 2);
// di will be a tuple<double,int, char>
auto di = make_tuple(2.5, 3, 'c');
cout << get<0>(di) // prints 2.5
cout << get<char>(di); // prints 'c' (C++14)
int three = get<1>(di);
```

Tuples

- Tuples are very useful for creating compound types on the fly
- We will implement an improved version of tuple from scratch in a few weeks

One annoyance with tuple vs pair

- pair<int, int> f() { return {1, 2}; // ok }
- tuple<int, int, char> f() { return {1, 2, 'u'}; // Error }
- We'll examine the reason for this and fix it in our own tuple class
- For more information, see Improving Pair and Tuple (revision 1) by Daniel Krugler
  - http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2013/n3739.html
  - Warning: This won't make sense until next quarter

Working with time

- Many programs work with time periods, but traditionally they just use integers or typedef to represent time
  - clock_t, time_t, dwMilliseconds, etc.
- This is very type-unsafe and error prone
  - For example, accidentally giving the number of milliseconds to a function that expects a number of seconds
  - Compiler won't even warn

durations

- C++ provides a `chrono::duration` type that represents an interval of time
- `duration` is actually a template class that indicates what the "clock tick" is
    - You can get the number of ticks with the `count()` method
- Usually, you don't use the template arguments directly (they are a little complicated), instead there are typedefs and literals for the different time units (C++14)
    ```
    using namespace std::literals::chrono_literals;
    using namespace std::chrono;
    auto threeSeconds = 3s;
    cout << threeSeconds.count() << " seconds" << endl;
    minutes minutesInTwoHours = 2h;
    // Casting milliseconds to hours is unsafe, so needs
    // an explicit cast
    hours h = duration_cast<hours>(123456ms);
    ```
- The literal suffixes leverage a feature called "user-defined literals"
    - Later, we will learn how to create our own literals this way

time_point

- Rather than a duration, a time_point represents a particular point in time

- Again, it's easier to use standard functions to create them

-
```
auto start=system_clock::now();
/* ... */
auto end = system_clock::now();
// elapsedTime will be a duration
auto elapsedTime = end - start;
cout << "Task took "
    << duration_cast<seconds>(elapsedTime).count()
    << " seconds" << endl;
```

Homework

Homework 9.1

- This problem consists of a series of types. Write a program that defines variables of each type set to some meaningful value (You are highly encouraged to check with a compiler). If the type is callable, the program should call it. Googling "c++ declarators" may help. Each one you get is worth 2 points.

- Example problem 1: `int *`
  - One possible answer:
    ```
    int *ip = new int;
    ```
  - Another possible answer
    ```
    int i = 5;
    int *ip = &i;
    ```

- Example problem 2: `int &`
  - One possible answer:
    ```
    int i = 5;
    int &ir(i);
    ```

HW 9.1 (cont)

- `int *`
- `int &`
- `double`
- `A *` (A is any appropriate class).
- `char const *`
- `char const &`
- `long[7]`
- `int **`
- `int *&`
- `float &`
- `int (*)()` (See http://www.newty.de/fpt/index.html)
- `int (*&)()`
- `char *(*)(char *, char *)`

HW 9.1 (cont)

- See http://www.informit.com/guides/content.aspx?g=cplusplus&seqNum=142 or the standard

- `int A::*`

- `int (A::*)(int *)`

- `int (A::**)(int *)`

- `int (A::*&)(int *)`

- `int (A::*)(double (*)(float &))`

- `void (*p[10]) (void (*)() );`

HW 9.2

- In slide 5, the function `f2` would be better with a different signature based on the best practices we've emphasized. What should it be?

- Unfortunately, even if you use `unique_ptr` and `shared_ptr` correctly, it is still possible to leak an object (i.e., the object is not deleted when you are done using it). Explain how this can happen

  - For extra credit, discuss how you might handle this