

C++


March 12, 2015

Mike Spertus

mike_spertus@symantec.com



New and Old Best Practices



Guard against multiple inclusion



- If the same header is included multiple times, you may get multiply-defined symbol errors
- The following preprocessor idiom prevents that from happening
 - Sometimes the preprocessor is helpful!
- ```
#ifndef FOO_H
define FOO_H
...
#endif
```

# Always put headers in a namespace



- Also use an “`#ifndef ...`” to guard against multiple inclusions
- ```
#ifndef FOO_H
#   define FOO_H
namespace mpcs51044 {
int f();
...
}
#endif
```

Never “use” a namespace in a header



- Leaks entire namespace to any file that includes the header.
- E.g., when in a header file, say
`using std::accumulate`
instead of
`using namespace std;`
or just explicitly call `std::accumulate`
without a `using` statement at all
- When in a “.cpp” file, choose whichever you prefer.

Prefer the C++ versions of standard C headers



- `#include <stdio.h> // Bad`
`#include <cstdio> // Better`
- The C versions will sort-of work, but the C++ versions will more properly define signatures, so overload resolution, type-checking, etc. will be more robust
- If you have a C header with no C++-specific version (e.g., `unistd.h`), then of course use the C version

Prefer C++-style casts to C style casts



- `A *a = (A *)&b; // bad`
- `A *a = dynamic_cast<A *>(&b);`

Prefer C++-style casts to C style casts -- Rationale



- Let's look at two cases where they differ

```
struct X {...};
struct Y {...};
X *xp = new X;
Y *yp = (X *) (xp); // Nonsense
yp = dynamic_cast<Y *>(xp); // 0 shows cast failed
```

```
struct Z : public X {...};
struct W : public X {...};
struct A : public Z, public W {...};
W *wp = new A;
X *xp = wp; // OK. Inheritance
A *ap = (A *)xp; // Oops! Points in middle of A
ap = dynamic_cast<A *>(xp); // Adjusts for
                           // multiple inheritance
```

- In both cases, C++-style casts are better when they disagree

Put const and volatile after type names



- “int const” is better and more consistent than “const int”
- Bjarne Stroustrup disagrees
- However, Dan Sachs’ ACCU “Truthiness” keynote argues this is the only rational conclusion one can reach, as it is both more logical and studies show that it leads to fewer bugs.

Use nullptr instead of 0 to indicate a null pointer



- C++ adds a new literal `nullptr` of type `nullptr_t` that represents (surprise) a null pointer. Automatically converts to pointer types (and bool)

```
void f(char *) { /* ... */ }  
void f(int) { /* ... */ }
```

```
f(0); // OK. Calls f(int)
```

```
f(nullptr); // OK. Calls f(char *)
```

- Always prefer the type-correct `nullptr` over the type-incorrect `0` or `NULL` to avoid calling the wrong function/method.

Define symmetric binary operators as global functions



- Don't use the member form of `operator+()`
 - Because both arguments should be treated the same
- However, do define `operator+=()` as a member
 - We don't want to `+=` to assign to a compiler-generated temporary

Think about types inferred by templates



- What does this print?

```
double dp[] = { 0.1, 0.2, 0.3 };  
cout << accumulate(dp, dp + 3, 0);
```

Think about types inferred by templates



- If you're accumulating doubles with `std::accumulate` use an initial value of `0.0` instead of `0`
 - Or you'll accumulate integers
- E.g.,

```
double dp[] = { 0.1, 0.2, 0.3 };  
cout << accumulate(dp, dp + 3, 0);
```

(surprisingly) prints 0

Beware of Dependent base classes



- What does the following print?

```
#include <iostream>
using namespace std;

int f() { return 0; }
template<class T>
struct C : public T {
    C() { cout << f() << endl; }
};
struct A {
    int f() { return 1; }
};
int main()
{
    C<A> c;
}
```

Dependent base classes: Surprising answer



- Microsoft Visual C++ prints 1
- g++ prints 0
- g++ is correct
- T is a “dependent base class”
 - A base class that depends on the template parameter
- Symbols are not looked up in dependent base classes, so templates are not surprised by unexpected inheritance

Correct use of dependent base classes



- To see symbols in a dependent base class, reference it explicitly:

```
template<class T>
struct C : public T {
    C() { cout << T::f() << endl; }
};
```

- Alternatively

```
template<class T>
struct C : public T {
    using T::f;
    C() { cout << f() << endl; }
};
```

- Tristan's choice

```
template<class T>
struct C : public T {
    C() { cout << this->f() << endl; }
};
```

- If you want the global symbol:

```
template<class T>
struct C : public T {
    C() { cout << ::f() << endl; }
};
```




Watch out for method hiding

```
struct B {  
    void f(bool i) { cout << "bool" << endl; }  
};  
  
struct D : public B {  
    // Fix with "using B::f"  
    void f(int b) { cout << "int" << endl; }  
};  
  
int main()  
{  
    D d;  
    d.f(true); // Prints "int"  
}
```

Use override and final to indicate intent



- ```
struct Base {
 virtual void func() = 0;
 virtual void misspelledFunc() = 0;
};
struct Derived : public Base {
 virtual void func() final {}
 // This will give a useful error
 // because we aren't actually
 // overriding
 virtual void misspelledFunc() override {}
};
struct MostDerived : public Derived {
 // Error! Can't override final
 virtual void func() { /*...*/ }
};
```
- This will catch a lot of “method hiding” errors

# Throw exceptions by value catch them by (const) &



- ```
struct MyException : public exception {
    MyException(string s)
        : myS("My "+s), exception(s) {}
    virtual char const *override what() {
        return myS.c_str();
    }
    string myS;
};

void f() {
    try {
        throw MyException("foo");
    } catch (exception e) { // Bad!
//} catch (exception const &e) { // Better
        cout << e.what(); // May crash due to slicing
    }
```

Never have a destructor throw an exception



- Does the following catch “In A” or “in f”?

```
struct A {  
    ~A() { throw runtime_error("In A"); }  
};  
void f()  
{  
    try {  
        A a;  
        throw runtime_error("in f");  
    } catch (exception const &) {  
    }  
}
```

- No good answer, so the runtime just calls `std::terminate` to end your program



Use const appropriately

- Const methods should be const
- Const & arguments should be const
- The “const” keyword should go after the type
- ```
class A {
 public:
 void f(int const &i) const;
};
```

# Use const appropriately-rationale



- Ignoring const is no longer an option
- ```
int seven() { return 7; }  
void pr_int(int &i) { cout << i; }  
void pr_int_const(int const &i) { cout << i; }  
pr_int(7); // Error  
pr_int(seven()); // Error on newer compilers  
pr_int_const(seven()); // OK
```
- Putting const on right prevents ambiguity
 - `const int *` looks like a constant “`int *`” but isn’t
 - `int const *` could only mean one thing
 - Studies show programmers make fewer mistakes with this rule



Don't slice objects

- D inherits from B
- D d;
B b = d; // Almost certainly wrong

Use virtual destructors when you inherit



- ```
class A {
 public:
 // virtual ~A() {}
};
class B : public A {
 public:
 ~B() { ... }
};
A *ap = new B;
delete ap; // Doesn't call B's dest
```





# Prefer templates to macros

- e.g., `min` should be a template but Microsoft Visual C++ defines it as a macro

# Don't make tricky assumptions about order of evaluation



```
struct S {
 S(int i) : a(i), b(i++) {
 f(i, i++) // Undefined behavior
 }
 int b;
 int a;
};
```

# Remember that primitive types have trivial constructors



```
void
f()
{
 int i;
 /* int i{}; // Fix with */
 cout << i; // i contains garbage
}
```

# Don't return a reference/pointer to a local variable



- ```
int &
f()
{
    int i = 3;
    return i; // Bad!
}
```

Best practice—Prefer range member functions to their single-element counterparts



- Item 5 of Meyer's Effective STL
- Given two vectors, v1 and v2, what's the easiest way to make v1's contents be the same as the second half of v2's?
 - Don't worry whether v2 has an odd number of elements



Worst (but common)

- `vector<Widget> v1, v2`
...
`for (vector<Widget>::const_iterator ci`
 `= v2.begin() + v2.size() / 2;`
 `ci != v2.end();`
 `++ci) {`
 `v1.push_back(*ci);`
}

Better



- `copy(v2.begin() + v2.size()/2,
v2.end(),
back_inserter(v1));`



Better yet

- ```
v1.resize(v2.size() - v2.size()/2);
 copy(v2.begin() + v2.size()/2,
 v2.end(),
 v1.begin());
```





# Even better

- `v1.insert`  
    `(v1.end(),`  
        `v2.begin() + v2.size() / 2,`  
        `v2.end());`

# Best



- `v1.assign(v2.begin() + v2.size()/2, v2.end());`

# Best Practice: Prefer `empty()` to `size() == 0`



- Suppose `l` is a `list<int>`
- Which is better?
  - `if (l.empty()) { ... }`
  - `if (l.size() == 0) { ... }`
- Prefer the `l.empty()`
- Calculating `size()` can take a long time
- Effective STL Item 4

# Recall the difference between virtual and non-virtual



- Review slides 32-35 of lecture 2
- This *will* be on the final

# Always use a smart pointer to manage the lifetime of an object



- `unique_ptr` if it has only one owner `shared_ptr` if it has multiple owners
- ```
Foo *fp(new Foo); // Bad  
unique_ptr<Foo> upfp(new Foo); // Good  
...  
delete fp; // May be missed if exception occurred
```
- More generally, use RAII to ensure resources get destroyed when they are no longer needed



Avoid using “new ...”

- The problem with saying “`new A()`” is that it returns an owning raw pointer to the new object, violating the preceding best practice
- Prefer using `make_unique` and `make_shared` **instead**
- `auto ap = make_unique<Foo>(); // Best`



Use RAII to manage locks

- Just like using smart pointers for objects, use a scoped locking class whose destructor releases the lock to make sure locks get released even when exceptions bypass normal control flow
 - Typically, this means to use the `std::lock_guard` class, like we do in the false sharing example
 - At work, I (Mike) just had a critical customer defect this week because manual unlocking code was bypassed by an exception.
 - Moral: Don't rely on manual unlocking code!



Lock ordering

- If you want to avoid deadlocks, you want to acquire locks in the same order!
 - Suppose thread 1 acquires lock A and then lock B
 - Suppose thread 2 acquires lock B and then lock A
 - There is a window where we could deadlock with thread 1 owning lock A and waiting for lock B while thread 2 owns lock B and is waiting for lock A forever
- The usual best practice is to document an order on your locks and always acquire them consistent with that order
- See
<http://www.ddj.com/hpc-high-performance-computing/204801163>

Memory model best practices



- Here are the takeaways
 - Try to avoid sharing data between threads except when necessary
 - When you share data between threads, always use locks or atomics to ensure both threads have a coherent view of the shared data
- A good reference
 - Boehm, Adve, “*You Don’t Know Jack about Shared Variables of Memory Models: Data Races are Evil*” Communications of the ACM 55, 2 Feb. 2012
 - <http://queue.acm.org/detail.cfm?id=2088916>



Make mutex members mutable

- ```
struct A {
 int f() const {
 lock_guard<mutex> lck(mtx);
 return i + j;
 }
 int i;
 int j;
 // So const methods can lock
 mutex mutable mtx;
};
```

# Final



- Open book
- Open notes
- You can look at posted sample files, lecture notes, your past HW submissions and the standard
  - You will definitely want to have ready access to the best practice list above
- Do not use a compiler
- Do not use any other resources or google for answers to questions